

# Computación Paralela: CUDA

JHON STEEVEN CABANILLA ALVARADO

MIGUEL CHAVEINTE GARCÍA

G01

# Declaración de bloques e hilos

```
1  //*****  
2  //                      Blocks & Threads                      //  
3  //*****  
4  // Declaración del numero de bloques y del numero de hilos por bloque  
5  int num_blocks;  
6  int num_blocks_reduction;  
7  int threads_per_block_reduction = 1024;  
8  int threads_per_block = 1024;  
9  
10 // Calculamos el numero de bloques en funcion del tamaño del problema  
11 num_blocks = num_searchers / threads_per_block;  
12 num_blocks_reduction = (rows * columns) / threads_per_block_reduction;  
13  
14 // Si el tamaño del vector no es multiplo del tamaño del bloque, lanzamos otro bloque completo (tendremos hilos sobrantes)  
15 if (num_searchers % threads_per_block != 0)  
16 {  
17     num_blocks++;  
18 }  
19  
20 if ((rows * columns) % threads_per_block_reduction != 0)  
21 {  
22     num_blocks_reduction++;  
23 }
```

# Lanzamientos de Kernels



```
1  /* 3.2. Terrain initialization */
2  // Inicializamos las matrices del device
3  computeInitiationMatrix<<<num_blocks_reduction, threads_per_block_reduction, 0, stream[0]>>>(dev_heights, dev_trails, rows * columns);
4  CHECK_CUDA_LAST();
5
6  /* 3.3. Searchers initialization */
7  searchersInitialization<<<num_blocks, threads_per_block, 0, stream[1]>>>(dev_searchers, dev_total_steps, num_searchers);
8  CHECK_CUDA_LAST();
9
10 /* 4. Compute searchers climbing trails */
11 computeSearchersTrails<<<num_blocks, threads_per_block, 0, stream[2]>>>(num_searchers, rows, columns, dev_searchers, dev_heights, dev_trails, x_min, x_max, y_min, y_max);
12 CHECK_CUDA_LAST();
13
14 /*4.1*/
15 computeFollows<<<num_blocks, threads_per_block, 0, stream[2]>>>(dev_searchers, dev_trails, num_searchers, columns);
16 CHECK_CUDA_LAST();
17
18 /* 5. Compute the leading follower of each searcher */
19 computeFollower<<<num_blocks, threads_per_block, 0, stream[2]>>>(dev_searchers, num_searchers);
20 CHECK_CUDA_LAST();
21
22 /* 6. Compute accumulated trail steps to each maximum */
23 computeAccumulatedTrail<<<num_blocks, threads_per_block, 0, stream[2]>>>(dev_searchers, dev_total_steps, num_searchers);
24 CHECK_CUDA_LAST();
```

# Cómputo estadístico

```
1 // Shared Memory
2 int shared_memory = threads_per_block * sizeof(int);
3 int global_reduction = threads_per_block_reduction * sizeof(int);
4
5
6 // Maximo accum steps
7 cudaMemsetAsync(max_steps, INT_MIN, sizeof(int));
8 cudaMalloc(&max_steps, sizeof(int)); CHECK_CUDA_LAST();
9 reductionMax<<<num_blocks, threads_per_block, shared_memory>>>(dev_total_steps, num_searchers, max_steps); CHECK_CUDA_LAST();
10 cudaMemcpyAsync(&max_accum_steps, max_steps, sizeof(int), cudaMemcpyDeviceToHost); CHECK_CUDA_LAST();
11
12 // Maxima altura ++ Total Heights
13 cudaMalloc(&max_altura, sizeof(int)); CHECK_CUDA_LAST();
14 cudaMalloc(&alturas, sizeof(unsigned long long int)); CHECK_CUDA_LAST();
15 reductionMax<<<num_blocks_reduction, threads_per_block_reduction, global_reduction>>>(dev_heights, rows * columns, max_altura); CHECK_CUDA_LAST();
16 reductionAddHeights<<<num_blocks_reduction, threads_per_block_reduction, global_reduction>>>(dev_heights, rows * columns, alturas); CHECK_CUDA_LAST();
17 cudaMemcpyAsync(&max_height, max_altura, sizeof(int), cudaMemcpyDeviceToHost); CHECK_CUDA_LAST();
18 cudaMemcpyAsync(&total_heights, alturas, sizeof(unsigned long long int), cudaMemcpyDeviceToHost); CHECK_CUDA_LAST();
19
20 // Num local max
21 cudaMalloc(&max_local, sizeof(int)); CHECK_CUDA_LAST();
22 computeLocalMax<<<num_blocks, threads_per_block, shared_memory>>>(dev_searchers, num_searchers, max_local); CHECK_CUDA_LAST();
23 cudaMemcpyAsync(&num_local_max, max_local, sizeof(int), cudaMemcpyDeviceToHost); CHECK_CUDA_LAST();
24
25 // Total tainted
26 cudaMalloc(&tainted_prueba, sizeof(int)); CHECK_CUDA_LAST();
27 reductionAdd<<<num_blocks_reduction, threads_per_block_reduction, global_reduction>>>(dev_trails, rows * columns, tainted_prueba); CHECK_CUDA_LAST();
28 cudaMemcpyAsync(&total_tainted, tainted_prueba, sizeof(int), cudaMemcpyDeviceToHost); CHECK_CUDA_LAST();
29
30 // Sincronizacion final de hilos
31 cudaDeviceSynchronize();
```

# Reducciones



```
1 // Load array values in the shared memory (0 if out of the array)
2 if (globalPos < size)
3 { // and array[globalPos] != INT_MIN
4     if (array[globalPos] != INT_MIN)
5     {
6         buffer[threadIdx.x] = array[globalPos];
7     }
8     else
9     {
10        buffer[threadIdx.x] = 0;
11    }
12 }
13 else
14     buffer[threadIdx.x] = 0;
15
16 // Wait for all the threads of the block to finish
17 __syncthreads();
18
```



```
1 // Reduction tree
2 for (int step = blockDim.x / 2; step >= 1; step /= 2)
3 {
4     if (threadIdx.x < step)
5         if (buffer[threadIdx.x] < buffer[threadIdx.x + step])
6             buffer[threadIdx.x] = buffer[threadIdx.x + step];
7     __syncthreads();
8 }
```

# Condiciones de carrera



```
1  /* Stop if searcher finds another trail */
2      int check;
3      // check = atomicAdd(&accessMat( tainted, pos_row, pos_col ), 1);
4      check = atomicCAS(&accessMat(trails, pos_row, pos_col), -1, search);
5
6      if (check != -1)
7      {
8          searchers[search].follows = check;
9          search_flag = 1;
10     }
```

# Ejemplo de kernel



```
1  __global__ void computeFollows(Searcher *searchers, int *trails, int num_searchers, int columns)
2  {
3
4      // Calculamos la posicion global de cada hilo
5      int gid = threadIdx.x + blockDim.x * blockIdx.x;
6
7      // Podemos tener más hilos de los necesarios. Por ello, añadimos este IF para que los hilos ociosos no entren en esta parte del código
8      if (gid < num_searchers)
9      {
10         int pos_row = searchers[gid].pos_row;
11         int pos_col = searchers[gid].pos_col;
12
13         searchers[gid].follows = accessMat(trails, pos_row, pos_col);
14     }
15 }
```

# Optimizaciones

- ▶ Eliminar matriz tainted
- ▶ Eliminar el atributo id del struct Searchers
- ▶ Declaración de variables para ahorrar llamadas a funciones
- ▶ Uso de Streams



# Streams



```
1  /* 3.2. Terrain initialization */
2  // Inicializamos las matrices del device
3  computeInitialitatonMatrix<<<num_blocks_reduction, threads_per_block_reduction, 0, stream[0]>>>(dev_heights, dev_trails, rows * columns);
4  CHECK_CUDA_LAST();
5
6  /* 3.3. Searchers initialization */
7  searchersInitialization<<<num_blocks, threads_per_block, 0, stream[1]>>>(dev_searchers, dev_total_steps, num_searchers);
8  CHECK_CUDA_LAST();
9
10 /* 4. Compute searchers climbing trails */
11 computeSearchersTrails<<<num_blocks, threads_per_block, 0, stream[2]>>>(num_searchers, rows, columns, dev_searchers, dev_heights, dev_trails, x_min, x_max, y_min, y_max);
12 CHECK_CUDA_LAST();
13
14 /*4.1*/
15 computeFollows<<<num_blocks, threads_per_block, 0, stream[2]>>>(dev_searchers, dev_trails, num_searchers, columns);
16 CHECK_CUDA_LAST();
17
18 /* 5. Compute the leading follower of each searcher */
19 computeFollower<<<num_blocks, threads_per_block, 0, stream[2]>>>(dev_searchers, num_searchers);
20 CHECK_CUDA_LAST();
21
22 /* 6. Compute accumulated trail steps to each maximum */
23 computeAccumulatedTrail<<<num_blocks, threads_per_block, 0, stream[2]>>>(dev_searchers, dev_total_steps, num_searchers);
24 CHECK_CUDA_LAST();
```

# Pinned Memory



```
1  //*****  
2      //                Pinned Memory                //  
3  //*****  
4  /*  
5      cudaMallocHost( (void **)&searchers, sizeof(Searcher) * num_searchers); CHECK_CUDA_LAST();  
6      cudaMallocHost( (void **)&total_steps, sizeof(int) * num_searchers); CHECK_CUDA_LAST();  
7  
8      cudaMallocHost( (void **)&heights, sizeof(int) * (size_t)rows * (size_t)columns); CHECK_CUDA_LAST();  
9      cudaMallocHost( (void **)&trails, sizeof(int) * (size_t)rows * (size_t)columns); CHECK_CUDA_LAST();  
10     cudaMallocHost( (void **)&tainted, sizeof(int) * (size_t)rows * (size_t)columns); CHECK_CUDA_LAST();  
11     */
```



```
1 // Copia de datos del Host al Device
2 //*****//
3 //          HostToDevice          //
4 //*****//
5
6 /*
7 cudaMemcpyAsync( dev_heights, heights, sizeof(int) * (size_t)rows * (size_t)columns, cudaMemcpyHostToDevice, stream[0]); CHECK_CUDA_LAST();
8 cudaMemcpyAsync( dev_trails, trails, sizeof(int) * (size_t)rows * (size_t)columns, cudaMemcpyHostToDevice, stream[1]); CHECK_CUDA_LAST();
9 cudaMemcpyAsync( dev_tainted, tainted, sizeof(int) * (size_t)rows * (size_t)columns, cudaMemcpyHostToDevice, stream[2]); CHECK_CUDA_LAST();
10 cudaMemcpyAsync( dev_searchers, searchers, sizeof(Searcher) * num_searchers, cudaMemcpyHostToDevice, stream[3]); CHECK_CUDA_LAST();
11 cudaMemcpyAsync( dev_total_steps, total_steps, sizeof(int) * num_searchers, cudaMemcpyHostToDevice, stream[4]); CHECK_CUDA_LAST();
12 */
13
14 //cudaMemcpyAsync(dev_heights, heights, sizeof(int) * (size_t)rows * (size_t)columns, cudaMemcpyHostToDevice);
15 //CHECK_CUDA_LAST();
16 //cudaMemcpyAsync(dev_trails, trails, sizeof(int) * (size_t)rows * (size_t)columns, cudaMemcpyHostToDevice);
17 //CHECK_CUDA_LAST();
18 cudaMemcpyAsync(dev_searchers, searchers, sizeof(Searcher) * num_searchers, cudaMemcpyHostToDevice);
19 CHECK_CUDA_LAST();
20 //cudaMemcpyAsync(dev_total_steps, total_steps, sizeof(int) * num_searchers, cudaMemcpyHostToDevice);
21 //CHECK_CUDA_LAST();
```